

```
# Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns #visualisation
import matplotlib.pyplot as plt #visualisation
%matplotlib inline
sns.set# Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns #visualisation
import matplotlib.pyplot as plt #visualisation
%matplotlib inline
sns.set(color_codes=True)
```

```
# Loading the CSV file into a pandas dataframe.
df = pd.read_csv("CARS.csv")
df.head(5)
```

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horsepower	MPG_City
0	Acura	MDX	SUV	Asia	All	\$36,945	\$33,337	3.5	6.0	265	1
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820	\$21,761	2.0	4.0	200	2
2	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990	\$24,647	2.4	4.0	200	2

```
# Removing irrelevant features
df = df.drop(['Model', 'DriveTrain', 'Invoice', 'Origin', 'Type'], axis=1)
df.head(5)
```

	Make	MSRP	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	Wheelbase	Length
0	Acura	\$36,945	3.5	6.0	265	17	23	4451	106	189
1	Acura	\$23,820	2.0	4.0	200	24	31	2778	101	172
2	Acura	\$26,990	2.4	4.0	200	22	29	3230	105	183
3	Acura	\$33,195	3.2	6.0	270	20	28	3575	108	186
4	Acura	\$43,755	3.5	6.0	225	18	24	3880	115	197

```
# To identify the type of data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 428 entries, 0 to 427
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Make            428 non-null    object
1   MSRP            428 non-null    object
2   EngineSize      428 non-null    float64
3   Cylinders       426 non-null    float64
4   Horsepower      428 non-null    int64
5   MPG_City        428 non-null    int64
6   MPG_Highway     428 non-null    int64
7   Weight          428 non-null    int64
8   Wheelbase       428 non-null    int64
9   Length          428 non-null    int64
dtypes: float64(2), int64(6), object(2)
memory usage: 33.6+ KB
```

```
# Getting the number of instances and features
df.shape
```

```
(428, 10)
```

```
# Getting the dimensions of the data frame
df.ndim
```

```
df = df.drop_duplicates(subset='MSRP', keep='first')
df.count()
```

```
Make          410
MSRP          410
EngineSize    410
Cylinders     408
Horsepower    410
MPG_City      410
MPG_Highway   410
Weight        410
Wheelbase     410
Length        410
dtype: int64
```

```
# Finding the null values
print(df.isnull().sum())
```

```
Make          0
MSRP          0
EngineSize     0
Cylinders      2
Horsepower     0
MPG_City       0
MPG_Highway    0
Weight         0
Wheelbase      0
Length         0
dtype: int64
```

```
df.Horsepower.max()
```


```
500
```

```
# Printing the null value rows
df[240:242]
```

	Make	MSRP	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	W
247	Mazda	\$25,700	1.3	NaN	197	18	25	3053	
248	Mazda	\$27,200	1.3	NaN	238	18	24	3029	

```
sns.FacetGrid(df, hue='MSRP', height=5).map(plt.bar, 'EngineSize', 'Horsepower').add_legend()
```





```
# Filling the rows with the mean of the column
```

```
val = df['Cylinders'].mean()
df['Cylinders'][247] = round(val)
val = df['Cylinders'].mean()
df['Cylinders'][248] = round(val)
```

```
<ipython-input-19-d9e0250049dd>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Cylinders'][248] = round(val)
```




```
# Removing the formatting
```

```
df['MSRP'] = [x.replace('$', '') for x in df['MSRP']]
df['MSRP'] = [x.replace(',', '') for x in df['MSRP']]
df['MSRP'] = pd.to_numeric(df['MSRP'], errors='coerce')
```

84165

```
sns.boxplot(x=df['MSRP'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f90de5d4df0>
```




```
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
```

```
IQR = Q3 - Q1
print(IQR)
```

```
MSRP      19086.50
EngineSize    1.55
Cylinders     2.00
Horsepower    85.00
MPG_City       4.00
MPG_Highway    5.00
Weight       872.25
Wheelbase      9.00
Length       16.00
dtype: float64
```

```
sns.boxplot(x=df['MSRP'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f90de44ebe0>
```




```
df.describe()
```

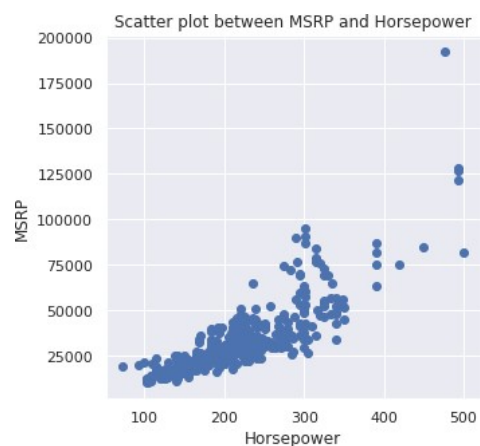
	MSRP	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	Wheelbas
count	410.000000	410.000000	410.000000	410.000000	410.000000	410.000000	410.000000	410.000000
mean	32919.021951	3.200976	5.824390	216.007317	20.039024	26.792683	3581.221951	108.16829
std	19628.241002	1.115264	1.556873	72.330689	5.280720	5.790004	766.396651	8.36563
min	10280.000000	1.300000	3.000000	73.000000	10.000000	12.000000	1850.000000	89.00000
25%	20324.750000	2.325000	4.000000	165.000000	17.000000	24.000000	3102.000000	103.00000
50%	27807.500000	3.000000	6.000000	210.000000	19.000000	26.000000	3476.000000	107.00000
75%	39411.250000	3.875000	6.000000	250.000000	21.000000	29.000000	3974.250000	112.00000
max	192465.000000	8.300000	12.000000	500.000000	60.000000	66.000000	7190.000000	144.00000

```
# Plotting a heat map
plt.figure(figsize=(10,5))
c= df.corr()
sns.heatmap(c,cmap="BrBG",annot=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f90de5d4fa0>




```
# Plotting a scatter plot
fig, ax = plt.subplots(figsize=(5,5))
ax.scatter(df['Horsepower'], df['MSRP'])
plt.title('Scatter plot between MSRP and Horsepower')
ax.set_xlabel('Horsepower')
ax.set_ylabel('MSRP')
plt.show()
```



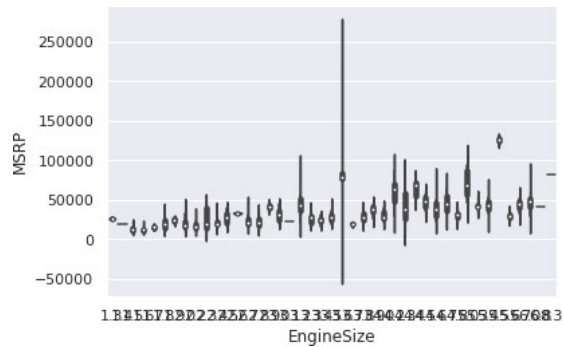
```
sns.jointplot(x='Horsepower', y='MSRP', data=df, height=10)
```

```
<seaborn.axisgrid.JointGrid at 0x7f90d9b91940>
```



```
sns.violinplot(x='EngineSize', y='MSRP', data=df, height=6)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f90d74f41c0>
```




```
sns.pairplot(df, hue='MSRP', height=2)
```

```
<seaborn.axisgrid.PairGrid at 0x7f90d69b7d30>
```



```
df.boxplot(by='MSRP', figsize=(12, 6))
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f90d6d84670>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d6257610>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d63053d0>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f90d7436430>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d74eb3d0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d6107520>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f90d613c490>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d750e670>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f90d62c8730>]],  
      dtype=object)
```



[Colab paid products](#) [Cancel contracts here](#)

✓ 12s completed at 10:41 PM

