

# Assignment 01

## Program for Matrix Addition :

```
%%cu
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <string.h>
#include <cuda.h>
#include <assert.h>

const int N = 4;
const int blocksize = 2;

__global__ void add_matrix_on_gpu( float* a, float *b, float *c, int N )
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    int index = i + j*N;
    if ( i < N && j < N )
        c[index] = a[index] + b[index];
}

void add_matrix_on_cpu(float *a, float *b, float *d)
{
    int i;
    for(i = 0; i < N*N; i++)
        d[i] = a[i]+b[i];
}

int main()
{
    printf("\n ***** CUDA Program for Matrix Addition *****\n");
    float *a = new float[N*N];
    float *b = new float[N*N];
    float *c = new float[N*N];
    float *d = new float[N*N];

    for ( int i = 0; i < N*N; ++i ) {
        a[i] = 1.0f; b[i] = 3.5f; }
}
```

```

printf("Matrix A:\n");
for(int i=0; i<N*N; i++)
{
    printf("\t%f",a[i]);
    if((i+1)%N==0)
        printf("\n");
}

printf("Matrix B:\n");
for(int i=0; i<N*N; i++)
{
    printf("\t%f",b[i]);
    if((i+1)%N==0)
        printf("\n");
}

    struct timeval  TimeValue_Start;
    struct timezone TimeZone_Start;

    struct timeval  TimeValue_Final;
    struct timezone TimeZone_Final;
    long            time_start, time_end;
    double          time_overhead;

float *ad, *bd, *cd;
const int size = N*N*sizeof(float);

cudaMalloc( (void**)&ad, size );
cudaMalloc( (void**)&bd, size );
cudaMalloc( (void**)&cd, size );

cudaMemcpy( ad, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( bd, b, size, cudaMemcpyHostToDevice );

dim3 dimBlock( blocksize, blocksize );
dim3 dimGrid( N/dimBlock.x, N/dimBlock.y );

gettimeofday(&TimeValue_Start, &TimeZone_Start);
add_matrix_on_gpu<<<dimGrid, dimBlock>>>( ad, bd, cd, N );
gettimeofday(&TimeValue_Final, &TimeZone_Final);

cudaMemcpy( c, cd, size, cudaMemcpyDeviceToHost );

```

```

add_matrix_on_cpu(a,b,d);

time_end = TimeValue_Final.tv_sec * 1000000 + TimeValue_Final.tv_u
sec;
time_start = TimeValue_Start.tv_sec * 1000000 + TimeValue_Start.tv
_usec;

time_overhead = (time_end - time_start)/1000000.0;

printf("result is:\n");
for(int i=0; i<N*N; i++)
{
    printf("\t%f%f",c[i],d[i]);
    if((i+1)%N==0)
        printf("\n");
}
for(int i=0; i<N*N; i++)
assert(c[i]==d[i]);

printf("\n\t\t Time in Seconds (T)           : %lf\n\n",time_overhea
d);

cudaFree( ad ); cudaFree( bd ); cudaFree( cd );
delete[] a; delete[] b; delete[] c, delete[] d;
return EXIT_SUCCESS;
}

```

## Output :

```

***** CUDA Program for Matrix Addition *****
Matrix A:
1.000000    1.000000    1.000000    1.000000
1.000000    1.000000    1.000000    1.000000
1.000000    1.000000    1.000000    1.000000
1.000000    1.000000    1.000000    1.000000
Matrix B:
3.500000    3.500000    3.500000    3.500000
3.500000    3.500000    3.500000    3.500000
3.500000    3.500000    3.500000    3.500000
3.500000    3.500000    3.500000    3.500000
result is:
4.5000004.500000    4.5000004.500000    4.5000004.500000    4.5000004.500000
4.5000004.500000    4.5000004.500000    4.5000004.500000    4.5000004.500000
4.5000004.500000    4.5000004.500000    4.5000004.500000    4.5000004.500000
4.5000004.500000    4.5000004.500000    4.5000004.500000    4.5000004.500000

Time in Seconds (T)           : 0.000012

```

## Program for transpose of matrix:

```
#include <cuda_runtime.h>
#include <iostream>

__global__ void transposeKernel(float* A, float* B, int m, int n) {
    int row = blockIdx.y*blockDim.y + threadIdx.y;
    int col = blockIdx.x*blockDim.x + threadIdx.x;
    if (row < n && col < m) {
        B[row*m + col] = A[col*n + row];
    }
}

int main() {
    int m = 3, n = 3;
    int size = m*n*sizeof(float);
    float* A, * B;
    cudaMalloc(&A, size);
    cudaMalloc(&B, size);
    // initialize A with sample values
    float A_host[9] = {1, 2, 3, 4, 5, 6, 7, 8, 9};
    cudaMemcpy(A, A_host, size, cudaMemcpyHostToDevice);
    dim3 block(3, 3);
    dim3 grid((m + block.x - 1) / block.x, (n + block.y - 1) / block.y);
    transposeKernel<<<grid, block>>>(A, B, m, n);
    cudaDeviceSynchronize();
    // print the input matrix A
    std::cout << "Matrix A:" << std::endl;
    for (int i = 0; i < 9; i++) {
        std::cout << A_host[i] << " ";
        if ((i + 1) % 3 == 0) std::cout << std::endl;
    }
    // print the transposed matrix B
    float B_host[9];
    cudaMemcpy(B_host, B, size, cudaMemcpyDeviceToHost);
    std::cout << "Matrix A transposed (A^T):" << std::endl;
    for (int i = 0; i < 9; i++) {
        std::cout << B_host[i] << " ";
        if ((i + 1) % 3 == 0) std::cout << std::endl;
    }
    cudaFree(A);
    cudaFree(B);
    return 0;
}
```

## Output:

```
➞ Matrix A:  
1 2 3  
4 5 6  
7 8 9  
Matrix A transposed (A^T):  
1 4 7  
2 5 8  
3 6 9
```

## Program for Matrix Multiplication :

```
%%cu
#include <cuda_runtime.h>
#include <iostream>

__global__ void matmulKernel(float* A, float* B, float* C) {
    int row = blockIdx.y*blockDim.y + threadIdx.y;
    int col = blockIdx.x*blockDim.x + threadIdx.x;
    if(row < 3 && col < 3) {
        float value = 0;
        for (int i = 0; i < 3; i++) {
            value += A[row*3 + i] * B[i*3 + col];
        }
        C[row*3 + col] = value;
    }
}

int main() {
    float* A, * B, * C;
    cudaMalloc(&A, 9 * sizeof(float));
    cudaMalloc(&B, 9 * sizeof(float));
    cudaMalloc(&C, 9 * sizeof(float));
    // initialize A and B with sample values
    float A_host[9] = {1, 2, 3, 4, 5, 6, 7, 8, 9};
    float B_host[9] = {9, 8, 7, 6, 5, 4, 3, 2, 1};
    cudaMemcpy(A, A_host, 9 * sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(B, B_host, 9 * sizeof(float), cudaMemcpyHostToDevice);
    dim3 block(3, 3);
    dim3 grid((3 + block.x - 1) / block.x, (3 + block.y - 1) / block.y);
    matmulKernel<<<grid, block>>>(A, B, C);
    cudaDeviceSynchronize();
    std::cout<<"\n ***** CUDA Program for Matrix Multiplicaion
***** \n";
    // print the input matrices A and B
    std::cout << "Matrix A:" << std::endl;
    for (int i = 0; i < 9; i++) {
        std::cout << A_host[i] << " ";
        if ((i + 1) % 3 == 0) std::cout << std::endl;
    }
    std::cout << "Matrix B:" << std::endl;
    for (int i = 0; i < 9; i++) {
```

```

        std::cout << B_host[i] << " ";
        if ((i + 1) % 3 == 0) std::cout << std::endl;
    }
    // print the result matrix C
    float C_host[9];
    cudaMemcpy(C_host, C, 9 * sizeof(float), cudaMemcpyDeviceToHost);
    std::cout << "Matrix C (A*B):" << std::endl;
    for (int i = 0; i < 9; i++) {
        std::cout << C_host[i] << " ";
        if ((i + 1) % 3 == 0) std::cout << std::endl;
    }
    cudaFree(A);
    cudaFree(B);
    cudaFree(C);
    return 0;
}

```

## Output:

-



```

***** CUDA Program for Matrix Multiplicaion *****
Matrix A:
1 2 3
4 5 6
7 8 9
Matrix B:
9 8 7
6 5 4
3 2 1
Matrix C (A*B):
30 24 18
84 69 54
138 114 90

```